Stat 9911
Principles of AI: LLMs
Large Language Model Architectures 03

Edgar Dobriban

Department of Statistics and Data Science, the Wharton School, University of Pennsylvania

January 30, 2025

Penn
UNIVERSITY of PENNSYLVANIA

Wharton
UNIVERSITY of PENNSYLVANIA

Plan

- We continue with the details of transformer architectures.

Table of Contents

Order and Permutation Equivariance

- Order matters! "The cat chased the mouse." vs "The mouse chased the cat."
- A key consideration in language models is how to take position information into account.
- Claim: Full attention is *permutation equivariant*. If we permute the order of the input E, the order of the output $\hat{E}$ is permuted accordingly.
- Now consider the reps e_c and e_m of cat and mouse in "The cat chased the mouse."
    1. Rep should capture that e_c is the attacker and e_m is the victim.
    2. However, if I permute the sentence to "The mouse chased the cat.", the reps e_c and e_m of cat and mouse stay exactly the same! (due to perm. equiv.)
    3. So, e_c and e_m cannot capture the relation between cat/mouse.

Attention is Permutation Equivariant

- Permute input embeddings $e_1, \ldots, e_T$ by $T \times T$ permutation matrix Π (i.e., $E \to \Pi E$)
- Recalling $Q = EW'_q$, $K = EW'_k$, $V = EW'_v$, $Z = QK^\top / \sqrt{d'}$, this leads to:

$$Q \to \Pi Q, \quad K \to \Pi K, \quad V \to \Pi V, \quad Z \to \Pi Z \Pi^\top.$$

- Row-softmax composed with Exp preserves permutation equiv.:
  - Elementwise exponentiation: $\exp(\Pi Z \Pi^\top) = \Pi \exp(Z) \Pi^\top$.
  - Division by row-sums: Let $g(Z) = \text{diag}(Z1)^{-1} Z$. Now:

$$\text{diag}(\Pi Z \Pi^\top 1) = \text{diag}(\Pi Z 1) = \Pi \, \text{diag}(Z1) \, \Pi^\top$$

    Hence,
$$\text{diag}(\Pi Z \Pi^\top 1)^{-1} = \Pi \, \text{diag}(Z1)^{-1} \, \Pi^\top$$

    Thus: $g(\Pi Z \Pi^\top) = \Pi \, \text{diag}(Z1)^{-1} \, \Pi^\top \cdot \Pi \, Z \, \Pi^\top = \Pi \, g(Z) \, \Pi^\top$.

- Therefore $\hat{E} = AV \mapsto \Pi A \Pi^\top \Pi V = \Pi \hat{E}$. Full attention is permutation equivariant.

Adding Positional Information

▶ Causal attn is not permutation equivariant, and specific output embeddings are in general not perm. invariant (despite some claims).

▶ However, explicitly knowing about position can still help: Regardless of what input is, you should always pay more attention to last few tokens

▶ Standard self-attention does not have this built in and needs to learn it, as it is equally "eager to pay attention" to any previous token.

▶ To address this, early works adds a positional encoding matrix Γ (a vector for each location), to the embedding matrix E.

▶ Γ can be:
   ▶ Fixed (e.g., Vaswani et al. (2017))
   ▶ Learned (e.g., Radford et al. (2018, 2019); Brown et al. (2020))

Issues with Absolute Positional Encoding

- Absolute positions may not be as meaningful as relative positions (e.g., just-preceding tokens are influential).
- Designing embeddings that depend only on relative position aims to address this.

RoPE: Rotary Positional Encoding (Su et al., 2024)

▶ For context length T, define key and query functions:

$$k, q : \mathbb{R}^d \times [T] \to \mathbb{C}^{d'}$$

such that for some $g : \mathbb{R}^d \times \mathbb{R}^d \times [T] \to \mathbb{R}$:

$$\langle k(e_i, i), q(e_j, j) \rangle = g(e_i, e_j, i - j).$$

▶ We want to find k, q such that the above holds for any $e = e_i, e' = e_j \in \mathbb{R}^d$ and $i, j = 1, \dots,$

▶ Complex numbers are used for convenience and we will eventually use reals. Need real-valued result for softmax probability.

▶ A particular solution is of the form (Su et al., 2024):

$$k(e, j) = q(e, j) = q(e)e^{2\pi i \theta \cdot j}.$$

▶ Interpret θ as frequency: $j - i = 1/\theta$ is period of attention

▶ Frequency can be coordinate-specific, and Su et al. (2024) suggest using $\theta_t = F^{-2\lfloor (t-1)/2 \rfloor / d}$, for coords $t = 1, \dots, d$, for some F

▶ Llama 3 uses $F = 500,000$ (Dubey et al., 2024); see also here

Solving for Relative Positional Embedding

▶ Start with a simpler problem, where e_i, e_j are fixed and $k = q$, so we have

$$\langle k(i), k(j) \rangle = g(i - j).$$

Allow T to be arbitrarily large, and focus on one coordinate at a time.

▶ Then, we need to solve the following problem: Find nonzero complex-valued sequences $(a_n)_{n \geq 1}$ such that there exists a complex-valued sequence $(b_n)_{n \in \mathbb{Z}}$ for which:

$$a_m \overline{a_n} = b_{m-n}, \quad \text{for all } m, n \geq 1.$$

▶ Write $a_m = r_m e^{i\nu_m}$ and $b_n = s_n e^{i\eta_n}$ in polar form; $r_m, s_n > 0$ are uniquely determined.

▶ Using this, deduce the conditions:

$$\begin{cases} r_m r_n = s_{m-n} \\ \nu_m - \nu_n - \eta_{m-n} \in \mathbb{Z} \end{cases}$$

for all m, n.

Solution: Magnitudes and Phases

- From $r_m r_n = s_{m-n}$:
  - Take $m = n$: $r_m^2 = s_0$, for all m.
  - Therefore, $r_m = s_0^{1/2}$ does not depend on m.
- From $\nu_m - \nu_n - \eta_{m-n} \in \mathbb{Z}$:
  - Take $m = n + 1$: $\nu_{n+1} - \nu_n - \eta_1 \in \mathbb{Z}$.
  - Therefore, $\nu_{n+1} - \nu_1 - n\eta_1 \in \mathbb{Z}$.
  - Fractional part of ν_n determines the solution.
  - Without loss of generality, take $\nu_{n+1} = \nu_1 + n\eta_1$.
- Solution
$$a_m = s_0^{1/2} e^{2\pi i (\nu_1 + m\eta_1)}$$

- Equivalently, $a_m = Ce^{2\pi i m\theta}$, where $C \in \mathbb{C}$ and $\theta \in \mathbb{R}$; i.e.,
  $k(j) = Ce^{2\pi i \theta \cdot j}$
- Vectors: each coord. as above, $k(j) = (C_1 e^{2\pi i \theta_1 \cdot j}, C_2 e^{2\pi i \theta_2 \cdot j}, \ldots)$.
  Then,
$$\langle k(i), k(j) \rangle = \sum_m \left(C_m e^{2\pi i \theta_m i} \right) \left(\overline{C_m} e^{-2\pi i \theta_m j} \right)$$
$$= \sum_m |C_m|^2 e^{2\pi i \theta_m (i-j)}.$$

Real-valued Solution

▶ How to obtain real-valued solutions?

▶ Observe that

$$\overline{\langle k(i), k(j) \rangle} = \sum_m |C_m|^2 e^{-2\pi i \theta_m (i-j)}.$$

▶ So, if $k(j) = (e^{2\pi i \theta \cdot j}, e^{-2\pi i \theta \cdot j})$, then

$$\langle k(i), k(j) \rangle = e^{2\pi i \theta(i-j)} + e^{-2\pi i \theta(i-j)}$$
$$= 2\mathrm{Re}(e^{2\pi i \theta(i-j)}) = 2\cos(2\pi \theta(i-j))$$

Final Solution

▶ Return to the original problem, $\langle k(e,i), q(e',j)\rangle = g(e,e',i-j)$.

▶ Any functions of the form

$$k(e,i) = k(e)e^{2\pi i\,\theta\cdot i}, \; q(e',j) = q(e')e^{2\pi i\,\theta\cdot j}$$

are solutions with $g(e,e',i-j) = \langle k(e), q(e')\rangle e^{2\pi i\,\theta\cdot(i-j)}$. [could even have coordinate-specific phases]

▶ If we choose the coordinates of k, q to come in conjugate pairs, obtain $g(e,e',i-j) = 2\mathrm{Re}[\langle k(e), q(e')\rangle e^{2\pi i\,\theta\cdot(i-j)}]$

▶ Implement it:
  ▶ For any embeddings e_i, e_j and associated key and query k_i, q_j, also include their conjugates, and encode position via $(k_i e^{2\pi i\,\theta\cdot i}, \bar{k}_i e^{-2\pi i\,\theta\cdot i})/\sqrt{2}$, $(q_j e^{2\pi i\,\theta\cdot j}, \bar{q}_j e^{-2\pi i\,\theta\cdot j})/\sqrt{2}$
  ▶ Or $\mathrm{Re}[\langle k_i, q_j\rangle e^{2\pi i\,\theta\cdot(i-j)}]$ in the attn map

Final Solution in Real-Valued Form

- Let $k_i = k_{i1} + i k_{i2}, \quad q_j = q_{j1} + i q_{j2}$.
- Rotation matrix

$$R_a = R(\theta, a) = \begin{bmatrix} \cos(2\pi\theta a) & -\sin(2\pi\theta a) \\ \sin(2\pi\theta a) & \cos(2\pi\theta a) \end{bmatrix}.$$

- Attention inner product is:

$$\left\langle R(\theta, i) \begin{bmatrix} k_{i1} \\ k_{i2} \end{bmatrix}, R(\theta, j) \begin{bmatrix} q_{j1} \\ q_{j2} \end{bmatrix} \right\rangle$$

Illustration of rotary position embedding

Figure: Su et al. (2024)

Attention with Linear Biases (ALiBi)

▶ ALiBi (Press et al., 2022) reduces attention to past tokens, with a slope term m:

$$\text{softmax}(q_j K^\top - m[j-1, \ldots, 2, 1, 0]).$$

▶ Uses $m = 2^{-8h/H}$, where H is the number of attention heads and $h \in [H]$ is the index of the head (so different heads have different effective scope).

▶ Do not use any positional embedding.

▶ Implement by adding the biases to head-specific mask matrices (mask size: $H \times T \times T$)

▶ Empirically, generalizes better to unseen sequence lengths.

Length Extrapolation for ALiBi

Figure: Press et al. (2022)

Aside: Perplexity

- Perplexity: for an LM $p : V^* \rightarrow [0, 1]$, the perplexity of a string x is

$$q(x) = 2^{-\sum_{t=1}^{|x|} \log_2 p(x_t | x_{1:t-1})/|x|} = p(x)^{-1/|x|} \geq 1$$

Also, $q(x) = 2^{\text{avg entropy of string } x \text{ under } p \text{ in bits/token}}$. Sometimes use $\log_2 q(x) = \log[1/p(x)]/|x|$

- A smaller perplexity means that the string is assigned a higher probability.

- Smaller is better: small perplexity on test data is considered to represent a higher quality LM.

- The perplexity of the LM on a dataset D is $Q = \mathbb{E}_{X \sim D} q(X)$

Table of Contents

Towards Longer Contexts

- The context length T is the number of tokens that can be input to the LM. Longer contexts mean it has the potential to handle larger tasks
- Key bottleneck: Standard attention has a quadratic $\Theta(T^2)$ memory and computational complexity
- Idea: simplify and reduce attention mechanism

Sparse Attention (Child et al., 2019)

- ▶ Only attend to a small number of tokens
    - ▶ e.g., previous c tokens
- ▶ Consider H heads, where the h-th head at the j-th position can attend to the subset $A_j^{(h)} \subset \{1, \ldots, j\}$. Can view connectivity pattern as a graph.
- ▶ Factorized attention: As we stack layers, we are able to attend to i at j if $k_1 \in A_j^{(h_1)}$, $k_2 \in A_{k_1}^{(h_2)}$, ... $i \in A_{k_{m-1}}^{(h_m)}$

# Sparse Attention (Child et al., 2019)

Figure: Full and factorized attention patterns from Child et al. (2019)

▶ Something similar is used in GPT-3 (Brown et al., 2020)

Multi- and Grouped Query Attention

▶ Multi-Query Attention (Shazeer, 2019): keys and values are shared across all attention heads.

▶ Grouped Query Attention (Ainslie et al., 2023) is a generalization that shares single key and value heads for groups query heads, interpolating between multi-head and multi-query attention.

▶ Used in Llama 3 (Dubey et al., 2024).

References

J. Ainslie, J. Lee-Thorp, M. de Jong, Y. Zemlyanskiy, F. Lebron, and S. Sanghai. Gqa: Training generalized multi-query transformer models from multi-head checkpoints. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 4895–4901, 2023.

T. Brown, B. Mann, N. Ryder, M. Subbiah, J. D. Kaplan, P. Dhariwal, A. Neelakantan, P. Shyam, G. Sastry, A. Askell, et al. Language models are few-shot learners. *Advances in neural information processing systems*, 33: 1877–1901, 2020.

R. Child, S. Gray, A. Radford, and I. Sutskever. Generating long sequences with sparse transformers. *arXiv preprint arXiv:1904.10509*, 2019.

A. Dubey, A. Jauhri, A. Pandey, A. Kadian, A. Al-Dahle, A. Letman, A. Mathur, A. Schelten, A. Yang, A. Fan, et al. The llama 3 herd of models. *arXiv preprint arXiv:2407.21783*, 2024.

O. Press, N. Smith, and M. Lewis. Train short, test long: Attention with linear biases enables input length extrapolation. In *International Conference on Learning Representations*, 2022. URL https://openreview.net/forum?id=R8sQPpGCv0.

A. Radford, K. Narasimhan, T. Salimans, and I. Sutskever. Improving language understanding by generative pre-training, 2018.

A. Radford, J. Wu, R. Child, D. Luan, D. Amodei, and I. Sutskever. Language models are unsupervised multitask learners, 2019.

N. Shazeer. Fast transformer decoding: One write-head is all you need. *arXiv preprint arXiv:1911.02150*, 2019.

References

J. Su, M. Ahmed, Y. Lu, S. Pan, W. Bo, and Y. Liu. Roformer: Enhanced transformer with rotary position embedding. *Neurocomputing*, 568:127063, 2024.

A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, L. Kaiser, and I. Polosukhin. Attention is all you need. In *Advances in Neural Information Processing Systems*, 2017. URL https://arxiv.org/pdf/1706.03762.pdf.