

Stat 9911

Principles of AI: LLMs

Large Language Model Architectures 03

Edgar Dobriban

Department of Statistics and Data Science, the Wharton School, University of Pennsylvania

February 2, 2025



Plan

- ▶ We continue with the details of transformer architectures.

Table of Contents

Representing Position

Towards Longer Contexts

Order and Permutation Equivariance

- ▶ Order matters! "The **cat** chased the **mouse**." vs "The **mouse** chased the **cat**."
- ▶ A key consideration in language models is how to take position information into account.
- ▶ Claim: Full attention is *permutation equivariant*. If we permute the order of the input E , the order of the output \hat{E} is permuted accordingly.
- ▶ Now consider the reps e_c and e_m of **cat** and **mouse** in "The **cat** chased the **mouse**."
 1. Rep should capture that e_c is the attacker and e_m is the victim.
 2. However, if I permute the sentence to "The **mouse** chased the **cat**.", the reps e_c and e_m of **cat** and **mouse** stay exactly the same! (due to perm. equiv.)
 3. So, e_c and e_m cannot capture the relation between **cat**/**mouse**.

Attention is Permutation Equivariant

- ▶ Permute input embeddings e_1, \dots, e_T by $T \times T$ permutation matrix Π (i.e., $E \rightarrow \Pi E$)
- ▶ Recalling $Q = EW'_q, K = EW'_k, V = EW'_v, Z = QK^\top / \sqrt{d'}$, this leads to:

$$Q \rightarrow \Pi Q, \quad K \rightarrow \Pi K, \quad V \rightarrow \Pi V, \quad Z \rightarrow \Pi Z \Pi^\top.$$

- ▶ Row-softmax composed with Exp preserves permutation equiv.:
 - ▶ Elementwise exponentiation: $\exp(\Pi Z \Pi^\top) = \Pi \exp(Z) \Pi^\top$.
 - ▶ Division by row-sums: Let $g(Z) = \text{diag}(Z1)^{-1} Z$. Now:

$$\text{diag}(\Pi Z \Pi^\top 1) = \text{diag}(\Pi Z 1) = \Pi \text{diag}(Z 1) \Pi^\top$$

Hence,


$$\text{diag}(\Pi Z \Pi^\top 1)^{-1} = \Pi \text{diag}(Z 1)^{-1} \Pi^\top$$

$$\text{Thus: } g(\Pi Z \Pi^\top) = \Pi \text{diag}(Z 1)^{-1} \Pi^\top \cdot \Pi Z \Pi^\top = \Pi g(Z) \Pi^\top.$$

- ▶ Therefore $\hat{E} = AV \mapsto \Pi A \Pi^\top \Pi V = \Pi \hat{E}$. Full attention is permutation equivariant.

Adding Positional Information

- ▶ Causal attn is not permutation equivariant, and specific output embeddings are in general not perm. invariant (despite some [claims](#)).
- ▶ However, explicitly knowing about position can still help: Regardless of what input is, you should always pay more attention to last few tokens
- ▶ Standard self-attention does not have this built in and needs to learn it, as it is equally "eager to pay attention" to any previous token.
- ▶ To address this, early works adds a positional encoding matrix Γ (a vector for each location), to the embedding matrix E .
- ▶ Γ can be:
 - ▶ Fixed (e.g., [Vaswani et al. \(2017\)](#))
 - ▶ Learned (e.g., [Radford et al. \(2018, 2019\)](#); [Brown et al. \(2020\)](#))



Issues with Absolute Positional Encoding

- ▶ Absolute positions may not be as meaningful as relative positions (e.g., just-preceding tokens are influential).
- ▶ Designing embeddings that depend only on **relative position** aims to address this.

RoPE: Rotary Positional Encoding (Su et al., 2024)

- ▶ For context length T , define key and query functions:

$$k, q : \mathbb{R}^d \times [T] \rightarrow \mathbb{C}^{d'}$$

such that for some $g : \mathbb{R}^d \times \mathbb{R}^d \times [T] \rightarrow \mathbb{R}$:

$$\langle k(e_i, i), q(e_j, j) \rangle_{\mathbb{C}} = g(e_i, e_j, i - j).$$

Here $\langle a, b \rangle_{\mathbb{C}} = \sum_{l=1}^d a_l \bar{b}_l$, where for a complex number $z = x + iy$, $\bar{z} = x - iy$ denotes its complex conjugate

- ▶ We want to find k, q such that the above holds for any $e = e_i, e' = e_j \in \mathbb{R}^d$ and $i, j = 1, \dots$,
- ▶ Complex numbers are used for convenience and we will eventually use reals. Need real-valued result for softmax probability.

RoPE: Rotary Positional Encoding (Su et al., 2024)

- ▶ A particular solution is of the form (Su et al., 2024):

$$k(e, j) = q(e, j) = q(e)e^{2\pi i \theta \cdot j}.$$

- ▶ Interpret θ as frequency: $j - i = 1/\theta$ is period of attention
- ▶ Frequency can be coordinate-specific, and Su et al. (2024) suggest using $\theta_t = F^{-2\lfloor (t-1)/2 \rfloor / d}$, for coords $t = 1, \dots, d$, for some F
- ▶ Llama 3 uses $F = 500,000$ (Dubey et al., 2024); see also [here](#)

Solving for Relative Positional Embedding

- ▶ Start with a simpler problem, where e_i, e_j are fixed and $k = q$, so we have

$$\langle k(i), k(j) \rangle_{\mathbb{C}} = g(i - j).$$

Allow T to be arbitrarily large, and focus on one coordinate at a time.

- ▶ Then, we need to solve the following problem: Find nonzero complex-valued sequences $(a_n)_{n \geq 1}$ such that there exists a complex-valued sequence $(b_n)_{n \in \mathbb{Z}}$ for which:

$$a_m \overline{a_n} = b_{m-n}, \quad \text{for all } m, n \geq 1.$$

- ▶ Write $a_m = r_m e^{i\nu_m}$ and $b_n = s_n e^{i\eta_n}$ in polar form; $r_m, s_n > 0$ are uniquely determined.
- ▶ Using this, deduce the conditions:

$$\begin{cases} r_m r_n = s_{m-n} \\ \nu_m - \nu_n - \eta_{m-n} \in \mathbb{Z} \end{cases}$$

for all m, n .

Solution: Magnitudes and Phases

- ▶ From $r_m r_n = s_{m-n}$:
 - ▶ Take $m = n$: $r_m^2 = s_0$, for all m .
 - ▶ Therefore, $r_m = s_0^{1/2}$ does not depend on m .
- ▶ From $\nu_m - \nu_n - \eta_{m-n} \in \mathbb{Z}$:
 - ▶ Take $m = n + 1$: $\nu_{n+1} - \nu_n - \eta_1 \in \mathbb{Z}$.
 - ▶ Therefore, $\nu_{n+1} - \nu_1 - n\eta_1 \in \mathbb{Z}$.
 - ▶ Fractional part of ν_n determines the solution.
 - ▶ Without loss of generality, take $\nu_{n+1} = \nu_1 + n\eta_1$.
- ▶ Solution

$$a_m = s_0^{1/2} e^{2\pi i(\nu_1 + m\eta_1)}$$

- ▶ Equivalently, $a_m = C e^{2\pi i m\theta}$, where $C \in \mathbb{C}$ and $\theta \in \mathbb{R}$; i.e.,
 $k(j) = C e^{2\pi i \theta \cdot j}$
- ▶ Vectors: each coord. as above, $k(j) = (C_1 e^{2\pi i \theta_1 \cdot j}, C_2 e^{2\pi i \theta_2 \cdot j}, \dots)$.
Then,

$$\begin{aligned}\langle k(i), k(j) \rangle_{\mathbb{C}} &= \sum_m (C_m e^{2\pi i \theta_m i}) (\overline{C_m} e^{-2\pi i \theta_m j}) \\ &= \sum_m |C_m|^2 e^{2\pi i \theta_m (i-j)}.\end{aligned}$$

Real-valued Solution

- ▶ How to obtain real-valued solutions?
- ▶ Observe that

$$\overline{\langle k(i), k(j) \rangle_{\mathbb{C}}} = \sum_m |C_m|^2 e^{-2\pi i \theta_m (i-j)}.$$

- ▶ So, if $k(j) = (e^{2\pi i \theta \cdot j}, e^{-2\pi i \theta \cdot j})$, then

$$\begin{aligned}\langle k(i), k(j) \rangle_{\mathbb{C}} &= e^{2\pi i \theta (i-j)} + e^{-2\pi i \theta (i-j)} \\ &= 2\operatorname{Re}(e^{2\pi i \theta (i-j)}) = 2\cos(2\pi \theta (i-j))\end{aligned}$$

Final Solution

- ▶ Return to the original problem, $\langle k(e, i), q(e', j) \rangle_{\mathbb{C}} = g(e, e', i - j)$.
- ▶ Any functions of the form

$$k(e, i) = k(e)e^{2\pi i \theta \cdot i}, \quad q(e', j) = q(e')e^{2\pi i \theta \cdot j}$$

are solutions with $g(e, e', i - j) = \langle k(e), q(e') \rangle_{\mathbb{C}} e^{2\pi i \theta \cdot (i - j)}$. [could even have coordinate-specific phases]

- ▶ If we choose the coordinates of k, q to come in conjugate pairs, obtain $g(e, e', i - j) = 2\text{Re}[\langle k(e), q(e') \rangle_{\mathbb{C}} e^{2\pi i \theta \cdot (i - j)}]$
- ▶ Implement it:
 - ▶ For any embeddings e_i, e_j and associated key and query k_i, q_j , also include their conjugates, and encode position via $(k_i e^{2\pi i \theta \cdot i}, \bar{k}_i e^{-2\pi i \theta \cdot i})/\sqrt{2}, (q_j e^{2\pi i \theta \cdot j}, \bar{q}_j e^{-2\pi i \theta \cdot j})/\sqrt{2}$
 - ▶ Or $\text{Re}[\langle k_i, q_j \rangle_{\mathbb{C}} e^{2\pi i \theta \cdot (i - j)}]$ in the attn map

Final Solution in Real-Valued Form

- ▶ Let $k_i = k_{i1} + ik_{i2}$, $q_j = q_{j1} + iq_{j2}$.
- ▶ Rotation matrix

$$R_a = R(\theta, a) = \begin{bmatrix} \cos(2\pi\theta a) & -\sin(2\pi\theta a) \\ \sin(2\pi\theta a) & \cos(2\pi\theta a) \end{bmatrix}.$$

- ▶ Attention inner product is:

$$\left\langle R(\theta, i) \begin{bmatrix} k_{i1} \\ k_{i2} \end{bmatrix}, R(\theta, j) \begin{bmatrix} q_{j1} \\ q_{j2} \end{bmatrix} \right\rangle_{\mathbb{R}},$$

where $\langle \cdot, \cdot \rangle_{\mathbb{R}}$ is the usual inner product on \mathbb{R} .

Rotary position embedding




Figure: Su et al. (2024)

Generalization of rotary position embeddings

- ▶ We are using the structure of the translation group acting on the domain of sequences.
- ▶ We end up with group representations $e \mapsto R(\theta, a)e$
- ▶ Generalization to arbitrary domains and groups acting on them has been attempted: [1](#), [2](#)

Attention with Linear Biases (ALiBi)

- ▶ ALiBi (Press et al., 2022) reduces attention to past tokens, with a slope term m :

$$\text{softmax}(q_j K^\top - m[j - 1, \dots, 2, 1, 0]).$$

- ▶ Uses $m = 2^{-8h/H}$, where H is the number of attention heads and $h \in [H]$ is the index of the head (so different heads have different effective scope).
- ▶ Do not use any positional embedding.
- ▶ Implement by adding the biases to head-specific mask matrices (mask size: $H \times T \times T$)
- ▶ Empirically, generalizes better to unseen sequence lengths.

Length Extrapolation for ALiBi




Figure: [Press et al. \(2022\)](#)

Aside: Perplexity

- ▶ Perplexity: for an LM $p : V^* \rightarrow [0, 1]$, the perplexity of a string x is

$$q(x) = 2^{-\sum_{t=1}^{|x|} \log_2 p(x_t | x_{1:t-1}) / |x|} = p(x)^{-1/|x|} \geq 1$$

Also, $q(x) = 2^{\text{avg entropy of string } x \text{ under } p \text{ in bits/token}}$. Sometimes use $\log_2 q(x) = \log[1/p(x)]/|x|$

- ▶ A smaller perplexity means that the string is assigned a higher probability.
- ▶ Smaller is better: small perplexity on test data is considered to represent a higher quality LM.
- ▶ The perplexity of the LM on a dataset D is $Q = \mathbb{E}_{X \sim D} q(X)$

Table of Contents

Representing Position

Towards Longer Contexts

Towards Longer Contexts

- ▶ The context length T is the number of tokens that can be input to the LM. Longer contexts mean it has the potential to handle larger tasks
- ▶ Key bottleneck: Standard attention has a quadratic $\Theta(T^2)$ memory and computational complexity
- ▶ Idea: simplify and reduce attention mechanism

Sparse Attention (Child et al., 2019)

- ▶ Only attend to a small number of tokens
 - ▶ e.g., previous c tokens
- ▶ Consider H heads, where the h -th head at the j -th position can attend to the subset $A_j^{(h)} \subset \{1, \dots, j\}$. Can view connectivity pattern as a graph.
- ▶ Factorized attention: As we stack layers, we are able to attend to i at j if $k_1 \in A_j^{(h_1)}$, $k_2 \in A_{k_1}^{(h_2)}$, ... $i \in A_{k_{m-1}}^{(h_m)}$, for some k_1, k_2, \dots and h_1, h_2, \dots

Sparse Attention (Child et al., 2019)




Figure: Full and factorized attention patterns from Child et al. (2019). "The top row indicates, for an example 6x6 image, which positions two attention heads receive as input when computing a given output. The bottom row shows the connectivity matrix (not to scale) between all such outputs (rows) and inputs (columns)."

- Something similar is used in GPT-3 (Brown et al., 2020)

Sparse Attention Example

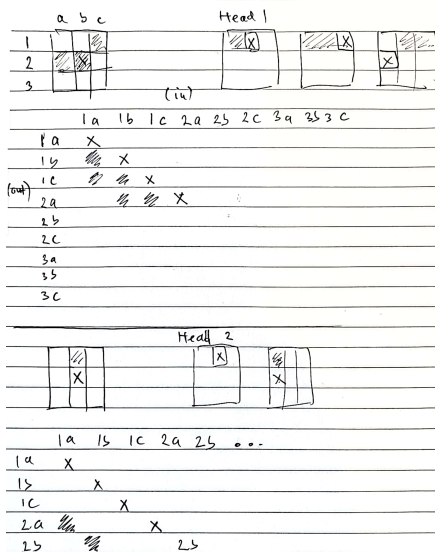



Figure: Full and factorized attention patterns from Child et al. (2019).

Multi- and Grouped Query Attention

- ▶ Multi-Query Attention ([Shazeer, 2019](#)): keys and values are shared across all attention heads.
- ▶ Grouped Query Attention ([Ainslie et al., 2023](#)) is a generalization that shares single key and value heads for groups query heads, interpolating between multi-head and multi-query attention.



- ▶ Used in Llama 3 ([Dubey et al., 2024](#)).

References

- J. Ainslie, J. Lee-Thorp, M. de Jong, Y. Zemlyanskiy, F. Lebron, and S. Sanghai. Gqa: Training generalized multi-query transformer models from multi-head checkpoints. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 4895–4901, 2023.
- T. Brown, B. Mann, N. Ryder, M. Subbiah, J. D. Kaplan, P. Dhariwal, A. Neelakantan, P. Shyam, G. Sastry, A. Askell, et al. Language models are few-shot learners. *Advances in neural information processing systems*, 33: 1877–1901, 2020.
- R. Child, S. Gray, A. Radford, and I. Sutskever. Generating long sequences with sparse transformers. *arXiv preprint arXiv:1904.10509*, 2019.
- A. Dubey, A. Jauhri, A. Pandey, A. Kadian, A. Al-Dahle, A. Letman, A. Mathur, A. Schelten, A. Yang, A. Fan, et al. The llama 3 herd of models. *arXiv preprint arXiv:2407.21783*, 2024.
- O. Press, N. Smith, and M. Lewis. Train short, test long: Attention with linear biases enables input length extrapolation. In *International Conference on Learning Representations*, 2022. URL <https://openreview.net/forum?id=R8sQPpGCv0>.
- A. Radford, K. Narasimhan, T. Salimans, and I. Sutskever. Improving language understanding by generative pre-training, 2018.
- A. Radford, J. Wu, R. Child, D. Luan, D. Amodei, and I. Sutskever. Language models are unsupervised multitask learners, 2019.
- N. Shazeer. Fast transformer decoding: One write-head is all you need. *arXiv preprint arXiv:1911.02150*, 2019.

References

- J. Su, M. Ahmed, Y. Lu, S. Pan, W. Bo, and Y. Liu. Roformer: Enhanced transformer with rotary position embedding. *Neurocomputing*, 568:127063, 2024.
- A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, L. Kaiser, and I. Polosukhin. Attention is all you need. In *Advances in Neural Information Processing Systems*, 2017. URL <https://arxiv.org/pdf/1706.03762.pdf>.